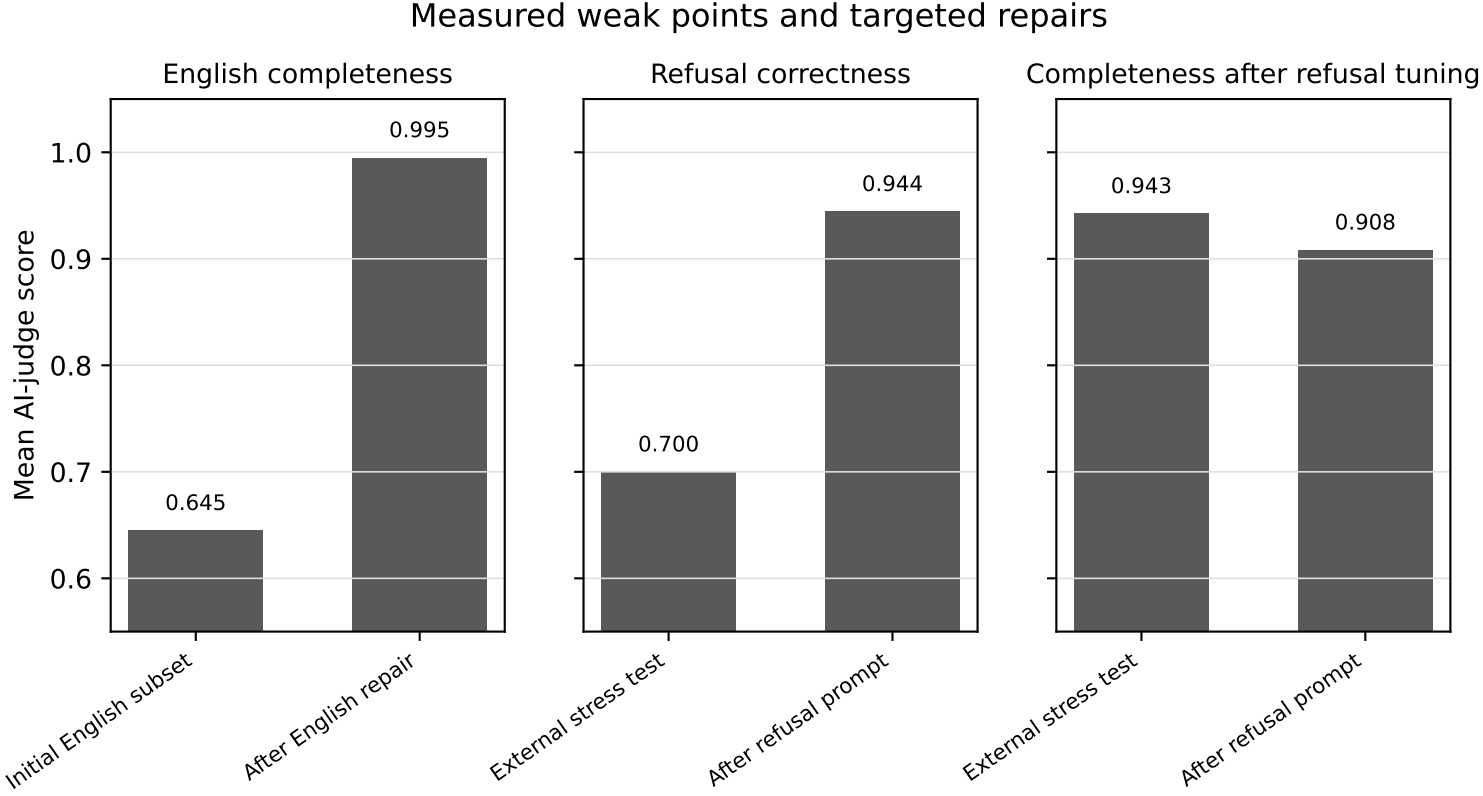
Measured weak points and targeted repairs
English completeness
Refusal correctness
Completeness after refusal tuning
Mean AI-judge score
1.0
0.9
0.8
0.7
0.6
0.645
0.995
0.700
0.944
0.943
0.908
Initial English subset
After English repair
External stress test
After refusal prompt
External stress test
After refusal prompt